

- 1-Import Libraries
 - 2-Import data
 - 3-Take care of na and missing data
 - 4-Except the dependent variable, convert all categorical variables to dummies
 - 5-Define X and scale it by using StandardScaler(it has a templete codes)
 - 6-Model the data (it has a templete codes)
 - 7-Evaluate a model using Silhouette_Score (it has a templete codes) (close to 1 is better clustering, close to 0 no clustering)
 - 8-Fit the optimal K on X-Scaled
 - 9- Create a new column for cluster's labels
 - 10- Groupby the clusters and show them in a graph
- Example is used from Build a K-means model as part of Google Advanced Data Analytics Professional Certificate

```
In [1]: # Import standard operational packages.
import numpy as np
import pandas as pd

# Important tools for modeling and evaluation.
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
from sklearn.preprocessing import StandardScaler

# Import visualization packages.
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: penguins = pd.read_csv(r"C:\Users\lasra\Downloads\penguins.csv")
```

```
In [4]: penguins.head(n=10)
```

```
Out[4]:
```

	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g	sex
0	Adelie	Torgersen	39.1	18.7	181.0	3750.0	male
1	Adelie	Torgersen	39.5	17.4	186.0	3800.0	female
2	Adelie	Torgersen	40.3	18.0	195.0	3250.0	female
3	Adelie	Torgersen	NaN	NaN	NaN	NaN	NaN
4	Adelie	Torgersen	36.7	19.3	193.0	3450.0	female
5	Adelie	Torgersen	39.3	20.6	190.0	3650.0	male
6	Adelie	Torgersen	38.9	17.8	181.0	3625.0	female
7	Adelie	Torgersen	39.2	19.6	195.0	4675.0	male
8	Adelie	Torgersen	34.1	18.1	193.0	3475.0	NaN
9	Adelie	Torgersen	42.0	20.2	190.0	4250.0	NaN

```
In [5]: penguins["species"].unique()

Out[5]: array(['Adelie', 'Chinstrap', 'Gentoo'], dtype=object)
```

```
In [7]: penguins["species"].value_counts(dropna=False)
```

```
Out[7]: Adelie      152
Gentoo    124
Chinstrap   68
Name: species, dtype: int64
```

There are three types of species. Note the Chinstrap species is less common than the other species.

This has a chance to affect K-means clustering as K-means performs best with similar sized groupings.

```
In [9]: penguins.isna().sum()
```

```
Out[9]: species      0
island      0
bill_length_mm    2
bill_depth_mm    2
flipper_length_mm 2
body_mass_g      2
sex          11
dtype: int64
```

```
In [10]: penguins_subset= penguins.dropna(axis=0).reset_index(drop=True)
```

```
In [11]: penguins_subset.isna().sum()
```

```
Out[11]: species      0
island      0
bill_length_mm    0
bill_depth_mm    0
flipper_length_mm 0
body_mass_g      0
sex            0
dtype: int64
```

```
In [12]: penguins_subset.head()
```

```
Out[12]:
```

	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g	sex
0	Adelie	Torgersen	39.1	18.7	181.0	3750.0	male
1	Adelie	Torgersen	39.5	17.4	186.0	3800.0	female
2	Adelie	Torgersen	40.3	18.0	195.0	3250.0	female
3	Adelie	Torgersen	36.7	19.3	193.0	3450.0	female
4	Adelie	Torgersen	39.3	20.6	190.0	3650.0	male

```
In [13]: penguins_subset['sex']=penguins_subset["sex"].str.upper()
```

```
In [14]: penguins_subset=pd.get_dummies(penguins_subset,drop_first=True, columns=["sex"])
```

```
In [16]: penguins_subset=penguins_subset.drop(['island'], axis=1)
```

```
In [18]: X=penguins_subset.drop(["species"], axis=1)
```

```
In [19]: X_scaled=StandardScaler().fit_transform(X)
```

```
In [21]: # Fit K-means and evaluate inertia for different values of k.

num_clusters = [i for i in range(2, 11)]

def kmeans_inertia(num_clusters, x_vals):
    """
    Accepts as arguments list of ints and data array.
    Fits a KMeans model where k = each value in the list of ints.
    Returns each k-value's inertia appended to a list.
    """
    inertia = []
    for num in num_clusters:
        kms = KMeans(n_clusters=num, random_state=42)
        kms.fit(x_vals)
        inertia.append(kms.inertia_)

    return inertia
```

```
In [22]: inertia = kmeans_inertia(num_clusters, X_scaled)
inertia
```

```
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=2.
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=2.
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=2.
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=2.
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=2.
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=2.
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less
```

```

chunks than available threads. You can avoid it by setting the environment variable O
MP_NUM_THREADS=2.
warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWa
rning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the va
lue of `n_init` explicitly to suppress the warning
warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWar
ning: KMeans is known to have a memory leak on Windows with MKL, when there are less
chunks than available threads. You can avoid it by setting the environment variable O
MP_NUM_THREADS=2.
warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWa
rning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the va
lue of `n_init` explicitly to suppress the warning
warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWar
ning: KMeans is known to have a memory leak on Windows with MKL, when there are less
chunks than available threads. You can avoid it by setting the environment variable O
MP_NUM_THREADS=2.
warnings.warn(

```

```

Out[22]: [885.6224143652249,
578.8284278107236,
386.14534424773296,
284.5464837898288,
217.92858573807678,
201.39287843423267,
185.461310432323,
173.45452114979847,
164.1200152026071]

```

```


In [23]: # Create a line plot.

```

```

plot = sns.lineplot(x=num_clusters, y=inertia, marker = 'o')
plot.set_xlabel("Number of clusters");
plot.set_ylabel("Inertia");

```



```
In [24]: # Evaluate silhouette score.
# Write a function to return a list of each k-value's score.

def kmeans_sil(num_clusters, x_vals):
    """
    Accepts as arguments list of ints and data array.
    Fits a KMeans model where k = each value in the list of ints.
    Calculates a silhouette score for each k value.
    Returns each k-value's silhouette score appended to a list.
    """
    sil_score = []
    for num in num_clusters:
        kms = KMeans(n_clusters=num, random_state=42)
        kms.fit(x_vals)
        sil_score.append(silhouette_score(x_vals, kms.labels_))

    return sil_score

sil_score = kmeans_sil(num_clusters, X_scaled)
sil_score
```

```
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=2.
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=2.
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=2.
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=2.
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=2.
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=2.
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less
```

```

chunks than available threads. You can avoid it by setting the environment variable O
MP_NUM_THREADS=2.
warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWa
rning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the va
lue of `n_init` explicitly to suppress the warning
warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWar
ning: KMeans is known to have a memory leak on Windows with MKL, when there are less
chunks than available threads. You can avoid it by setting the environment variable O
MP_NUM_THREADS=2.
warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWa
rning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the va
lue of `n_init` explicitly to suppress the warning
warnings.warn(
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWar
ning: KMeans is known to have a memory leak on Windows with MKL, when there are less
chunks than available threads. You can avoid it by setting the environment variable O
MP_NUM_THREADS=2.
warnings.warn(

```

```

Out[24]: [0.44398088353055243,
          0.45101024097188364,
          0.5080140996630784,
          0.519998574860868,
          0.5263224884981607,
          0.47774022332151733,
          0.4221920732643224,
          0.36062890821417276,
          0.3617250563420018]

```

```


In [25]: # Create a line plot.

```

```

plot = sns.lineplot(x=num_clusters, y=sil_score, marker = 'o')
plot.set_xlabel("# of clusters");
plot.set_ylabel("Silhouette Score");

```

```
In [26]: # Fit a 6-cluster model.
```

```
kmeans6 = KMeans(n_clusters=6, random_state=42)
kmeans6.fit(X_scaled)
```

C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

```
warnings.warn(
```

```
C:\Users\lasra\anaconda3-2\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=2.
```

```
warnings.warn(
```

Out[26]:

```
KMeans(n_clusters=6, random_state=42)
```

```
In [27]: # Print unique labels.
```

```
print('Unique labels:', np.unique(kmeans6.labels_))
```

```
Unique labels: [0 1 2 3 4 5]
```

```
In [28]: # Create a new column `cluster`.
```

```
penguins_subset['cluster'] = kmeans6.labels_  
penguins_subset.head()
```

Out[28]:

	species	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g	sex_MALE	cluster
0	Adelie	39.1	18.7	181.0	3750.0	1	2
1	Adelie	39.5	17.4	186.0	3800.0	0	1
2	Adelie	40.3	18.0	195.0	3250.0	0	1
3	Adelie	36.7	19.3	193.0	3450.0	0	1
4	Adelie	39.3	20.6	190.0	3650.0	1	2

In [29]: *# Verify if any `cluster` can be differentiated by `species`.*

```
penguins_subset.groupby(by=['cluster', 'species']).size()
```

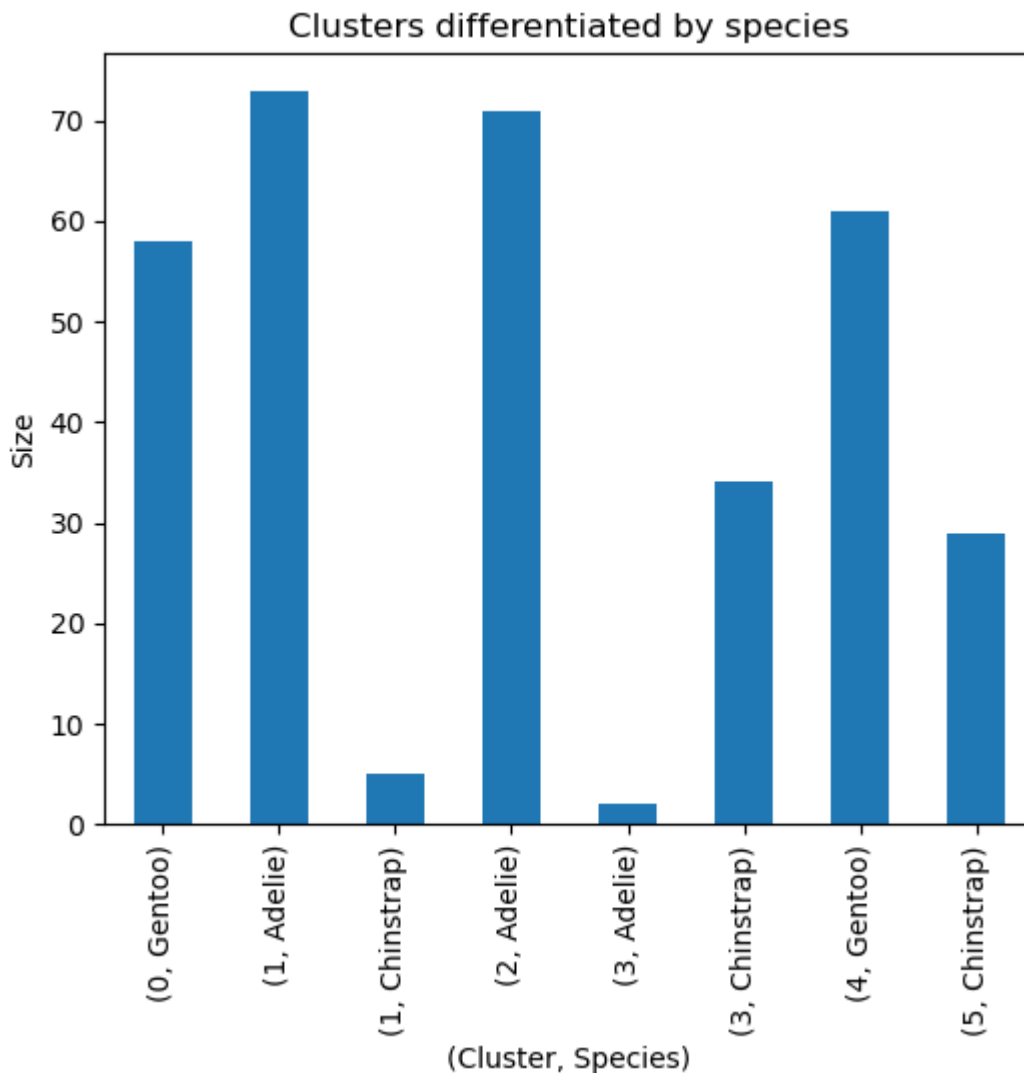
Out[29]:

cluster	species	
0	Gentoo	58
1	Adelie	73
	Chinstrap	5
2	Adelie	71
3	Adelie	2
	Chinstrap	34
4	Gentoo	61
5	Chinstrap	29

dtype: int64

In [30]:

```
penguins_subset.groupby(by=['cluster', 'species']).size().plot.bar(title='Clusters dif',
                                                                    figsize=(6, 5),
                                                                    ylabel='Size',
                                                                    xlabel='(Cluster, s
```



```
In [31]: # Verify if each `cluster` can be differentiated by `species` AND `sex_MALE`.
penguins_subset.groupby(by=['cluster', 'species', 'sex_MALE']).size().sort_values(ascending=True)
```

```
Out[31]:
```

cluster	species	sex_MALE	size
1	Adelie	0	73
2	Adelie	1	71
4	Gentoo	1	61
0	Gentoo	0	58
3	Chinstrap	1	34
5	Chinstrap	0	29
1	Chinstrap	0	5
3	Adelie	1	2

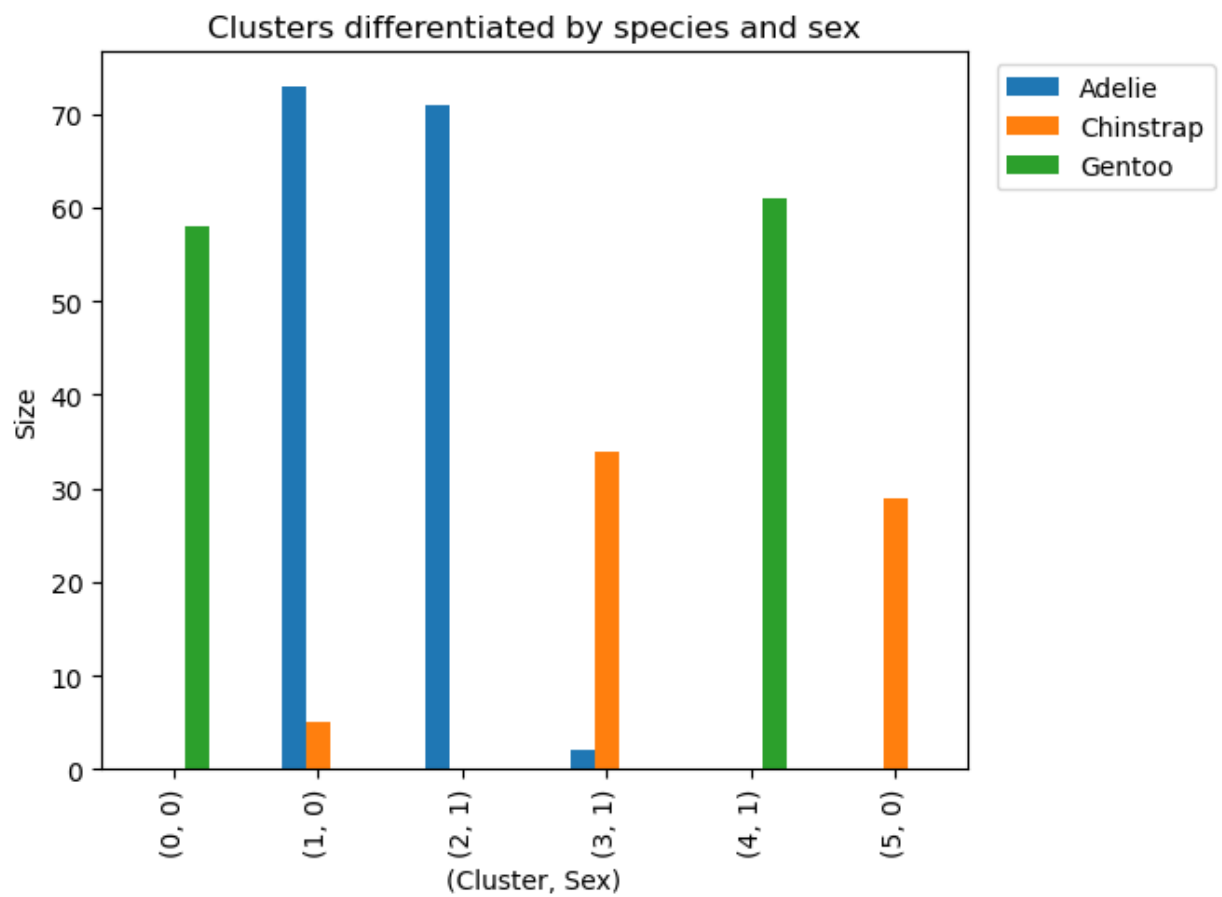
dtype: int64

Even though clusters 1 and 3 weren't all one species or sex, the groupby indicates that the algorithm produced clusters mostly differentiated by species and sex.

```
In [34]: penguins_subset.groupby(by=['cluster', 'species', 'sex_MALE']).size().\
unstack(level = 'species', fill_value=0).plot.bar(title='Clusters differentiated by species')

plt.legend(bbox_to_anchor=(1.3, 1.0))
```

Out[34]: <matplotlib.legend.Legend at 0x1d0328aa8c0>



In []: